

# *evaluating matches between data sets in R*

*alexis dhembe, lesley university library*



# *the problem*

At the Lesley library, we have several projects and recurring tasks that require comparing data between our ILS and other systems, such as student data from the bursar's office, or holdings data from our knowledge base. Because the two information sources were not always perfectly matched, previous workflows involved a lot of manual review in order to account for discrepancies between the two sources.



# *the goal*



I wanted to find a more efficient way to compare data sets, but I knew that we would have outliers that required some individual review. My goal was to find a way to compare the majority of records computationally, so staff only had to manually review records that actually warranted it.

Not all of the data sets had matching key identifiers I could use, so I also needed to account for a little fuzziness.

# two use cases

01

## *For Patron Data*

Each semester, we get a list of students who intend to graduate at the next conferral date. We wanted a better way to compare that to our patron data from the ILS to determine which students still had items checked out. We could match these accounts based on an institutional ID number, but I wanted to confirm that we didn't assign an ID number to the wrong patron at registration, or use an in-use ID as a "generic" stand-in. To typo is to be human, after all.

02

## *For Holdings Data*

We had a mystery package in our knowledge base that looked like it was meant to represent our print journal holdings, but details on the titles (such as ISSNs) were sparse. I wanted to compare the package to our current ILS holdings, but I needed to differentiate when something wasn't an exact match because of the difference between "Journal of..." and "Journal for..." or the difference between "i-D" and "Ms."

# enter: R's *stringdist* package

I've been using R (especially the *tidyverse*) to do a lot of the data wrangling that results from having information spread across several different systems, each with their own quirks. I knew I wanted to be able to compare inexact matches, and I wanted to be able to 'rank' those matches, so I could extract the more questionable ones for review.

R's *stringdist* package ended up being a perfect fit for my needs. *stringdist* provides fuzzy string matching capabilities based on various string distance measures. The basic *stringdist* function calculates the distance between two strings, e.g. the difference between "ABC" and "ACAB."

*stringsim* takes this distance, divides it by the maximum possible distance, and subtracts it from 1. This produces a value where 1 is perfect similarity, and 0 is completely dissimilar—precisely the kind of 'score' I could use to evaluate and rank matches.



# evaluating student data

I wanted to review records where the ID# matched between the bursar's records and ours, but the name differed. I created an additional key by combining last and first name into a single string. I joined the two tables on ID#, then used *stringsim* to create a new column comparing the name-key from the bursar records with ours. In order to err on the side of caution, given that this was a new workflow, everything with a similarity below 0.9 was flagged for additional review by library staff.

```
> grad_patrons %>% filter(SIMILARITY<0.9)
# A tibble: 79 x 11
```

	ID	NAMEKEY.x	LAST.x	FIRST.x	GRAD_DATE	NAMEKEY.y	LAST.y	FIRST.y	GROUP	ITEMS_OUT	SIMILARITY
		<dbl> <chr>	<chr>	<chr>	<date>	<chr>	<chr>	<chr>	<chr>	<chr>	<dbl>
1	800371	WARDJOHN	WARD	JOHN	2021-09-12	MARDJOHN	MARD	JOHN	UNDERGRAD	4	0.875
2	800653	LONGJENNIFER	LONG	JENNIFER	2021-09-12	YOUNGJENNIFER	YOUNG	JENNIFER	STAFF	2	0.846
3	800394	LAWRENCELAURA	LAWRENCE	LAURA	2021-05-16	LAWRENCETARA	LAWRENCE	TARA	PHD	4	0.846
4	800808	WATKINSMICHAEL	WATKINS	MICHAEL	2021-05-16	MATKINSMICHELE	MATKINS	MICHELE	STAFF	4	0.786
5	800958	CHENTARA	CHEN	TARA	2021-05-16	CHENDANA	CHEN	DANA	GRAD	4	0.75
6	800978	ROYMELINDA	ROY	MELINDA	2021-09-12	LEONMELINDA	LEON	MELINDA	PHD	1	0.727
7	800478	RASMUSSENMONICA	RASMUSSEN	MONICA	2021-05-16	RASWUSSENROBIN	RASWUSSEN	ROBIN	UNDERGRAD	3	0.667
8	800412	VELASQUEZROBYN	VELASQUEZ	ROBYN	2021-05-16	VELASQUEZPATRICK	VELASQUEZ	PATRICK	PHD	4	0.562
9	800733	DEANRANDALL	DEAN	RANDALL	2021-09-12	FORDRANDAL	FORD	RANDAL	UNDERGRAD	3	0.545
10	800323	ROBERSONKRISTINA	ROBERSON	KRISTINA	2021-09-12	ROBERSONANDREW	ROBERSON	ANDREW	STAFF	4	0.5

Note: this example uses 100% fake data created with Python's Faker package, not our real patron data!





# *student data (cont'd.)*

Our previous workflow for graduating patrons involved having staff and student workers look up graduates by name and review their accounts. By adding the additional step of evaluating account names, I was able to create a process that was much faster, but still had a safety mechanism for discrepancies (which was also reassuring to our stakeholders). We're able to identify ~93% of graduating patrons automatically, saving significant staff time.

# *comparing holdings data*

08



This method of evaluating string matches has also been useful for analyzing holdings data, particularly when one set is missing information or incomplete. In this particular case, we had a package in our knowledge base that seemed meant to represent our print holdings, for inclusion on our journals list. I wanted to compare it to our current print holdings (retrieved from our ILS) to see if it was still accurate and update coverage information. I first tried comparing titles by ISSN, and found matches in print holdings for 204 of the 469 titles from the Knowledge Base Mystery Package. This left 265 mystery titles remaining, which far exceeded my personal pain threshold for dragging spreadsheet cells around.



```
# A tibble: 182 x 5
  title.x          title.y
  <chr>          <chr>
1 action in teacher education
2 adta newsletter
3 adult learning
4 advances in mind body medicine
5 afterimage
6 american art
7 american art therapy association newsletter
8 american arts quarterly
9 american journal of dance therapy
10 american scientist
11 americas
12 archives of pediatrics adolescent medicine
13 art bulletin
14 art calendar
15 art journal
```

I decided to try fuzzy matching based on title, as I had little other data from the KB package to work with. I knew the *fuzzyjoin* package would be my best bet, using *stringdist\_inner\_join* to join the KB package to our catalog holdings on title. I cleaned extra punctuation and spaces and converted to lowercase to ease things along. This was fairly successful at first glance.

So far, so good, right? (Unlike the patron data shown earlier, this is our real holdings data; no use faking what's in the OPAC already)



holdings data (cont'd.)

I knew we had some very short titles, like the previously mentioned "i-D" and "Ms," where two edits could mean the difference between two totally different titles. I used *stringsim* to compare the matched titles from the KB and the print catalog, which allowed me to filter out those more tenuous matches. This isn't a huge dataset, but I've been able to apply this to other collection-management projects already.

# A tibble: 182 x 3			
	title.x	title.y	sim
	<chr>	<chr>	<dbl>
1	ms	id	0
2	ms	man	0.333
3	mome	bomb	0.5
4	wired	world	0.6
5	i d	id	0.667
6	spring	print	0.667
7	change	exchange	0.75
8	ilwc journal	iawm journal	0.833
9	photographis	photograph	0.833
10	viewcamera	view camera	0.909
11	lürzers intl archive	lurzers intl archive	0.95
12	brain mind common sense	brainmind common sense	0.957
13	advances in mind body medicine	advances in mindbody medicine	0.967
14	poliéster pintura y no pintura	poliester pintura y no pintura	0.967
15	action in teacher education	action in teacher education	1

# *fin!*

11

That's about it! This is a pretty simple presentation on a pretty simple function, but it's helped me tackle a lot of long-standing data management projects by filtering out what actually needs a more thorough review. It's also been helpful in demonstrating to library stakeholders that we can strike a balance between computational matching and manual review, which has saved a lot of staff-time and tedious clicking. It's been very useful for me, so I hope this can be useful to someone else as well!

***let's talk!***

github: [https://github.com/alxsdhm/c4l2021\\_poster](https://github.com/alxsdhm/c4l2021_poster)

tweets: @alxsdhm

discord: alxs#4443

